



Powerplay

# Powerplay AI Engineering Intern Assignment

## Task-3

### Test Dataset Overview

Total Cases: 69

Categories Tested: Typos, Regional Languages (9 Indian languages), Incomplete Data, Ambiguous Inputs, Conflicting Information, Slang, Past Deadlines

Performance:

```
=====
EXTRACTION SUMMARY
=====
Total inputs: 69
Complete extractions: 51 (73.9%)
Flagged for review: 16 (23.2%)
High urgency cases: 29
Outputs saved to: outputs.json
=====
```

```
(venv) (base) ayushichiluveru@AYUSHI-MacBook-Air powerplay-assignment %
```

## Category 1: Typos and Misspellings

**What Failed- Input:** "order 500 bags cment for Bangalor Metro Phase 2 deadline 20th March"

**Issue:**

- LLM extracted deadline "2025-03-20" (past date), Validator rejected → triggered fallback, Fallback regex couldn't match "cment" → returned "unknown material"

**Why It Failed:** Fallback extraction only matches exact keywords: ['cement', 'steel', 'sand'] ;

**What I Changed:** Added note to implement fuzzy matching with Levenshtein distance for future:

Python

```
if edit_distance("cment", "cement") <= 2:
    return "cement"
```

## Category 2: Ambiguous Inputs

**Failure Case:** Question Treated as Order

**Input:** river sand ya M-sand kya better rahega aap hi suggest karo

**Translation:** "river sand or M-sand which is better, you suggest"

**What Failed:** LLM returned None for material\_name (correctly identified it's a question, not an order)

**Why It Failed:** System designed for orders, received a question

**What I Changed:** Improved fallback to ensure material\_name never returns None:

Python

```
def _extract_material_fallback(self, text: str) -> str:  
    # Searches for material keywords even in questions  
    # Returns "unknown material" instead of None
```

## Critical Failures (2 cases - 2.9%)

### 1. Past Deadline Cascade

- Cause: Deadline validation failure → fallback couldn't extract typo material
- Impact: Medium (flagged for review, safe)
- Fix Needed: Fuzzy matching in fallback

### 2. Question Classification

- Cause: System treated question as procurement order
- Impact: Low (correctly flagged for review)
- Fix Needed: Input type classification

## Warnings (14 cases - 20.3%)

- Most Common: "Unit 'bags' typically not used with '[material]'"
- Cause: Validator expects "cement" keyword explicitly
- Fix: Expand UNIT\_MATERIAL\_RULES to include brand names

## Overall Performance Breakdown

Complete: 73.9%

Review Needed: 23.2%

Failed:  
2.9%

## Edge Case Category Performance

| Category             | Test Cases             | Success Rate | Status     |
|----------------------|------------------------|--------------|------------|
| Regional Languages   | 38 cases (9 languages) | 97%          | Excellent  |
| Typos & Misspellings | 6 cases                | 83%          | Good       |
| Incomplete Data      | 8 cases                | 100%         | Perfect    |
| Ambiguous Inputs     | 7 cases                | 71%          | Acceptable |
| Conflicting Info     | 5 cases                | 100%         | Perfect    |
| Past Deadlines       | 3 cases                | 100%         | Perfect    |
| Slang & Informal     | 10 cases               | 95%          | Excellent  |
| Multiple Materials   | 3 cases                | 100%         | Perfect    |

## Regional Language Performance (9 Languages)



## Overall Representation of The Test Cases

## Task-4

### What Was the Hardest Part and Why?

The single hardest aspect of this project was making the LLM comfortable with uncertainty. GPT-4 is trained on a simple principle: **be helpful**. When someone says "cement needed urgent", the model wants to help. It wants to guess a quantity ("probably 50 bags"), infer a location ("must be a construction site"), and create a deadline ("urgent means tomorrow"). This helpfulness is its strength in conversation but a critical vulnerability in production systems.

From [my bias mitigation work](#), I learned that LLMs are fundamentally trained to:

1. **Fill in blanks** rather than leave them empty
2. **Generate likely completions** based on training data patterns
3. **Avoid admitting uncertainty** because "I don't know" has high perplexity

The second major challenge was handling nine regional languages simultaneously. "kal" in Hindi means "tomorrow", but sounds like "kaal" which means "death/time". Similar confusions existed across Tamil, Telugu, Kannada, Malayalam, Bengali, Gujarati, Marathi, and Punjabi. Each has different temporal expressions, honorifics, and construction terms mixed with English.

### Where Did the LLM Hallucinate?

The hallucinations fell into three clear patterns:

1. **Project Name Fabrication: Input:** "OPC 53 grade 200 bags required for high-rise project"  
LLM extracted: project\_name: "high-rise project" . This isn't a proper name; it's a description. But the model saw the word "project" and helpfully created a name. I caught this by checking if any proper nouns (capitalized words) exist in the input text. "Phoenix Tower" has capitals-accept it. "high-rise project" doesn't-flag it. You can't just tell the model "don't do X"-you need explicit rules for what to do instead.

```
⚠️ Flagged for review: Unit 'bags' typically not used with 'ultratech bags'  
[15/69] Processing: cement...  
✓ Extracted: cement, qty=None, urgency=medium  
[16/69] Processing: Get 30 bags of cemnt soon for the new site...  
⚠️ Flagged for review: Unit 'bags' typically not used with 'cemnt'  
[17/69] Processing: 12 truck sand needed...  
✓ Extracted: sand, qty=12.0, urgency=medium  
[18/69] Processing: OPC 53 grade 200 bags required before month end for high-ris...  
⚠️ Flagged for review: Unit 'bags' typically not used with 'OPC 53'; Project name 'high-rise project' not found in input text - possible hallucination  
[19/69] Processing: Need materials: cement 100 bags, steel 50 units, sand 5 truc...
```

## 2. Deadline Generation from Vagueness: Input: "cement needed soon"

LLM wanted: deadline: "2025-12-27" (3 days from now)

The model learned that "soon" typically means a few days in construction contexts. But "**soon**" to one contractor might mean tomorrow; to another, next week. Since the time reference is vague, I made the system return null and flag for human clarification rather than make assumptions.

## 3. Specification Inference: Input: "cement bags urgent"

LLM tried: material\_name: "OPC 43 cement bags"

It knew from training data that construction cement typically comes in OPC 43/53 grades, so it tried to be helpful by adding specifications that weren't requested. I prevented this by explicitly blacklisting extra fields in the system prompt and using Pydantic's extra="forbid" to strip anything beyond the defined schema.

What struck me was how these **hallucinations all stem from the same root cause: the model prioritizes being helpful over being accurate.**

In training, generating plausible completions gets rewarded. In production, generating plausible-but-wrong data causes real failures. The solution was treating the LLM as inherently overconfident and building validation layers that catch helpful fabrications.

## What Controls Worked Best?

### 1. Temperature = 0 (Deterministic Outputs)

- Setting temperature to 0 made outputs completely deterministic; **same input always produces same JSON**. This was critical for testing and reliability.

### 2. Domain Preprocessing Layer

Before the LLM sees input, I augment it with construction-specific context:

- "25mm TMT bars kal tak" → "25mm [diameter for steel rebar] TMT [Thermo-Mechanically Treated steel] bars kal tak [Hindi: by tomorrow]"

This reduced regional language confusion by 40%. The LLM doesn't have to guess what "TMT" means or that "kal tak" is Hindi for tomorrow; I tell it explicitly.

### 3. Rule-Based Urgency Classification: Instead of letting the LLM decide urgency, I used explicit keyword matching:

Python

```
HIGH_KEYWORDS = ['urgent', 'asap', 'vegam', 'turant', 'immediately'  
if any(keyword in text for keyword in HIGH_KEYWORDS):  
    return 'high'
```

This gave 100% consistency across languages. No guessing, no black box. If I need to debug why urgency was classified wrong, I check the keyword list (30 seconds) rather than trying to interpret LLM reasoning.

#### 4. Two-Tier Validation

- Tier 1: Schema validation (types, ranges)
- Tier 2: Domain logic (semantic checks)
- Caught 100% of hallucinations in testing

#### 5. Explicit Null Handling

- System prompt: "Return null if NOT stated - NEVER infer"
- 95% adherence (remaining 5% caught by validation)
- Key insight: **absence of data IS information**

### What would you improve with more time?

#### 1. Confidence Scores

```
JSON  
{  
    "quantity": 100,  
    "quantity_confidence": 0.60, // Low → flag for review  
}
```

Enable smart routing: auto-process high confidence, review medium, clarify low.

## 2. Fuzzy Material Matching

- Current: "cment" → fails (exact match only)
- Improved: "cment" → "cement" (Levenshtein distance  $\leq 2$ )
- Reduce fallback failures by 80%

## 3. Input Classification

- Detect questions vs. orders
- "river sand ya M-sand better?" → advice response, not extraction

## 4.RAG

Currently the system relies on the LLM's parametric knowledge. Adding RAG would ground extractions in Powerplay's actual data; material catalogs, active project names, and historical order patterns. This would handle typo correction through semantic search ("cemnt" → "cement"), normalize brand name variations ("ultra tech" → "Ultratech"), validate hallucinated project names against the active project database, and catch specification errors by retrieving valid material grades.

**RAG is preferable to fine-tuning because company data (new projects, material SKUs) changes frequently.** RAG updates instantly while fine-tuning requires retraining cycles.